

Database Management System (DBMS)

— Vishvadeep Yothi playlist

Definition of DBMS

DBMS is a system designed to store, retrieve, update and manage data.

Advantages of DBMS

• Centralized data storage and retrieval.

• Consistency and integrity of data.

• Data independence and portability.

• Data redundancy and inconsistency reduction.

• Data security and access control.

• Scalability and performance optimization.

• Simplified data management and maintenance.

• Standardized interfaces and protocols.

• Improved data sharing and collaboration.

Database :- The collection of data, usually referred to as the database, contains information relevant to an enterprise.

Database Management System :- The database management system (DBMS) is a collection of interrelated data and a set of programs to access those data.

↳ info retrieval,
insertion, deletion,
update.

Advantages of DBMS

→ Providing a way to store & retrieve database info that is both convenient & efficient.

Goal of DBMS

→ Ensuring safety of information

(ACID properties)

(atomicity, consistency, integrity, durability)

pass swab no required
(ACID)

Disadvantages of file system

→ Data Redundancy and Inconsistency

→ Difficulty in Accessing Data

→ Data isolation

→ Integrity problems (applying constraints)

→ Atomicity problem

→ Concurrent - Access Anomalies

→ Security problems

- swab to wait

- lock disappears

- memory leak

- don't know

Instance - The collection of information stored in the database at a particular moment is called an instance of database.

Schema - The overall design of the database is called database schema.

Database Languages

Data Definition Language (DDL)

(schema related changes are done using DDL)

Data Manipulation Language (DML)

(data related changes are done using DML)

Procedural DMLs

requires user to specify what data is needed how to get that data

Non-procedural (Declarative-DMLs)

requires user to specify what data is needed without specifying how to get those data

View of data -

1. Physical level -

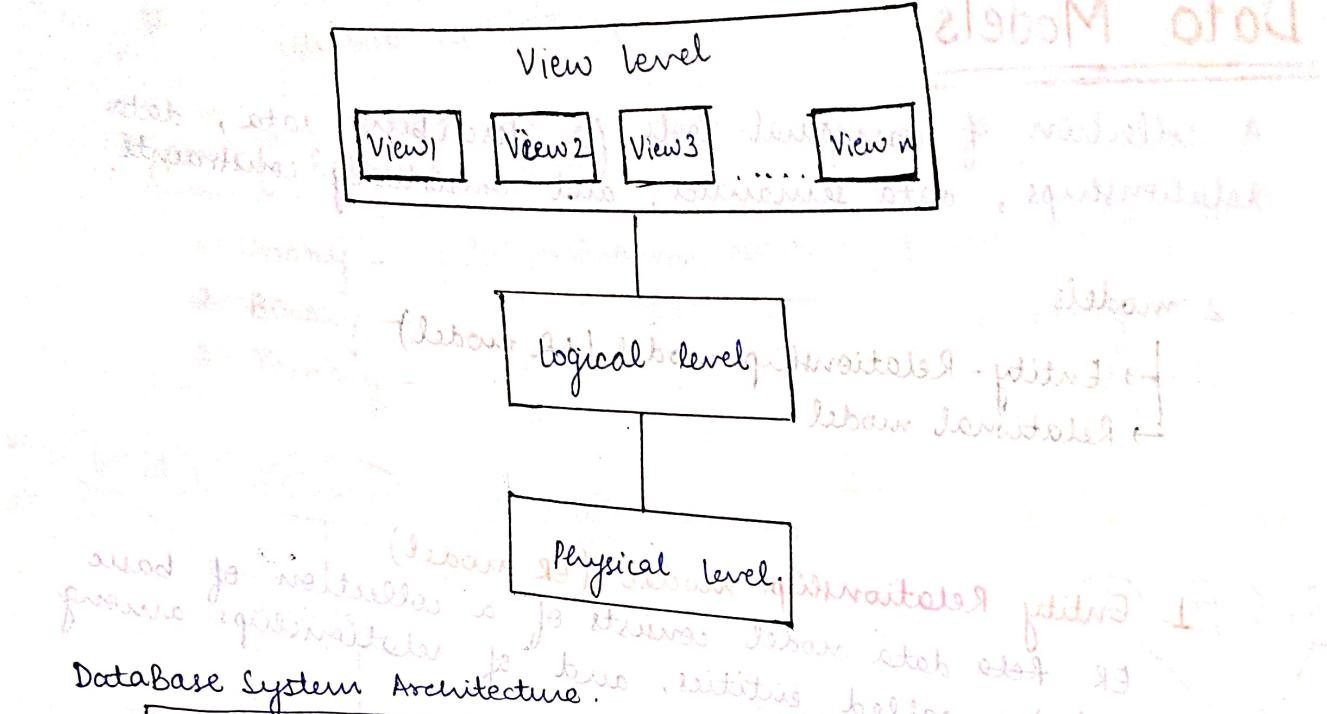
how data is actually stored in physical memory.

2. Logical level -

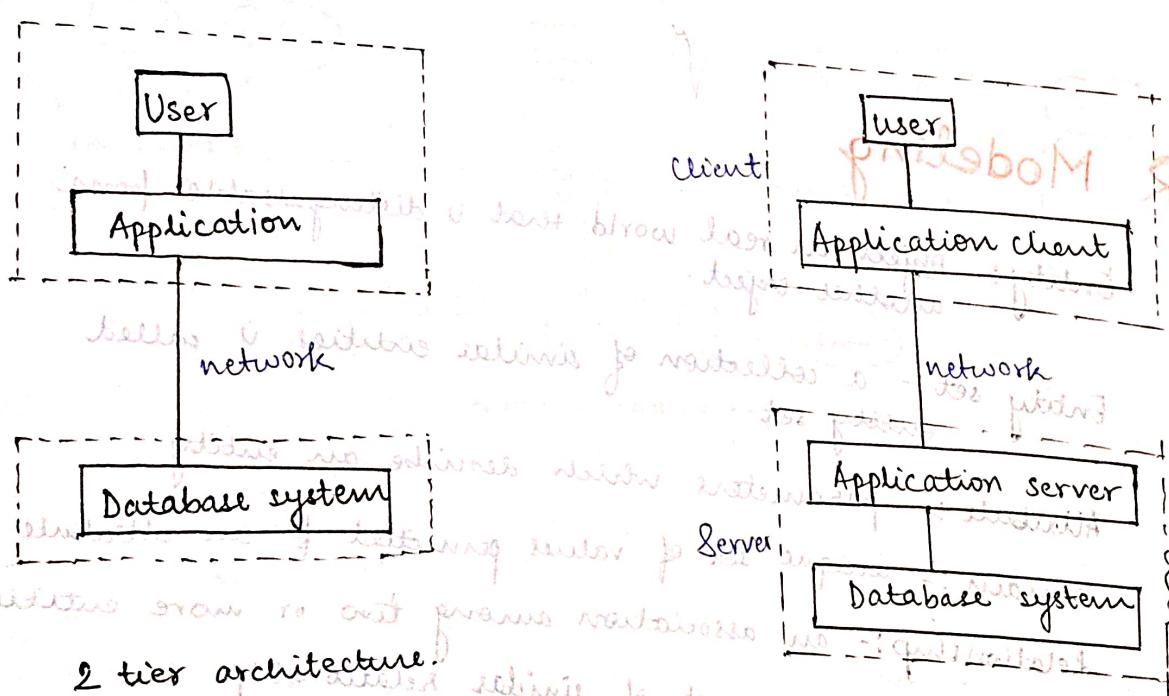
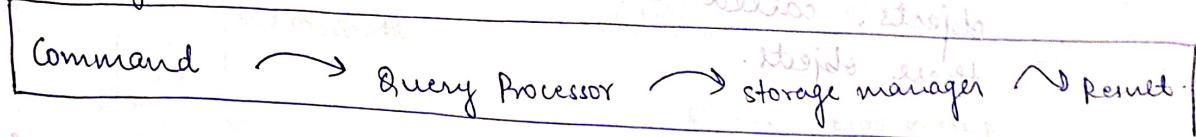
tables, at this level, we see data distributed in the form of tables in database.

3. View level - defines which part of database can be viewed by a person.

manager can view all the tables
employee can view only his record.



Database System Architecture



Data Models

A collection of conceptual tools for describing data, data relationships, data semantics, and consistency constraints.

2 models

→ Entity-Relationship model (ER model)

→ Relational model

1. Entity Relationship model (ER model)

ER data model consists of a collection of basic objects, called entities, and of relationships among these objects.

2. Relational model

uses a collection of tables to represent both data & relationships.

E-R Modeling

Entity :- object in real world that is distinguishable from another object.

Entity set :- a collection of similar entities is called entity set.

Attribute :- parameters which describe an entity

Domain :- unique set of values permitted for an attribute

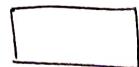
Relationship :- an association among two or more entities

Relationship-set - a set of similar relationships

Key - an attribute or a set of attributes whose values uniquely identify an entity in a set.

Descriptive attribute - attribute of a relationship.

Entity set →



Relationship →



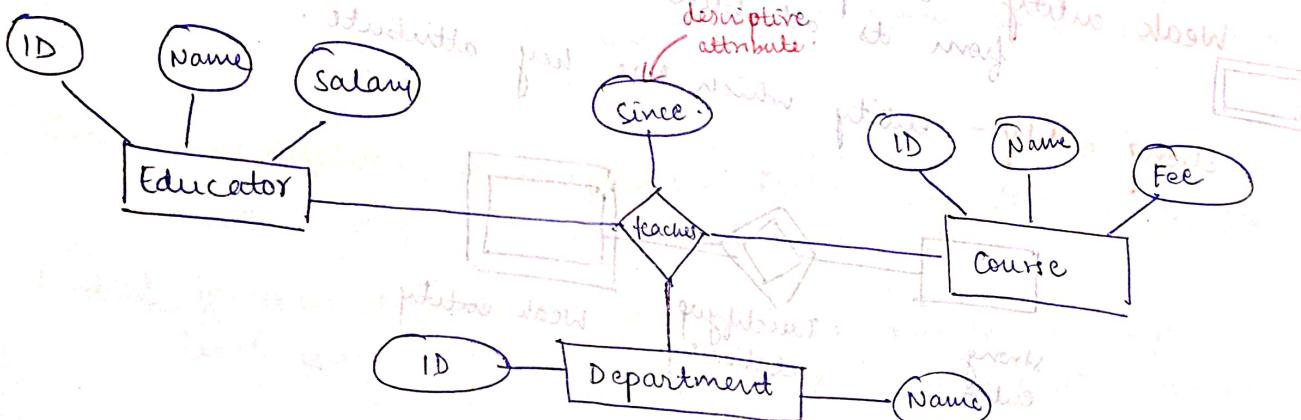
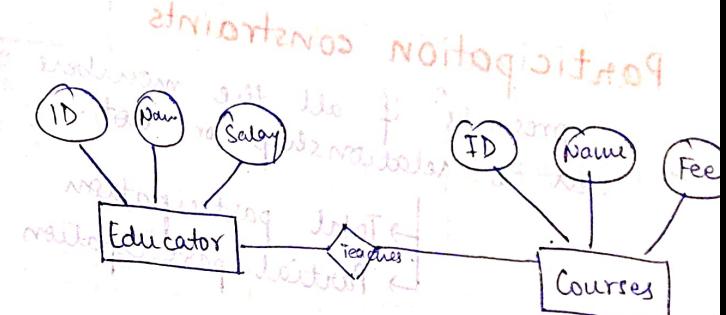
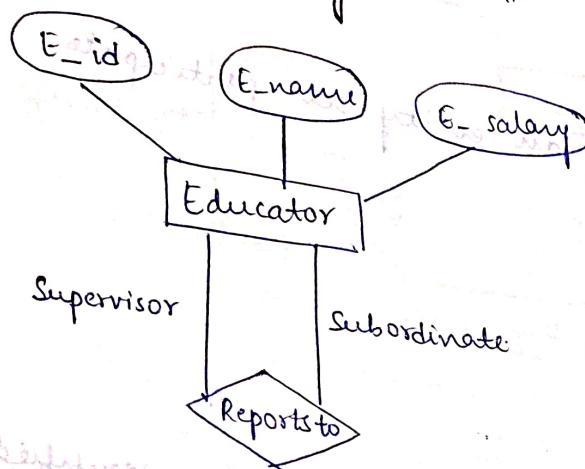
UNANIMOUS AGREEMENT

attribute →



Types of Relationships -

1. Unary - Relationship b/w entities of same entity set.
2. Binary - Relationship b/w entities of 2 entity sets.
3. Ternary - Relationship b/w entities of 3 entity sets.



Mapping Cardinality

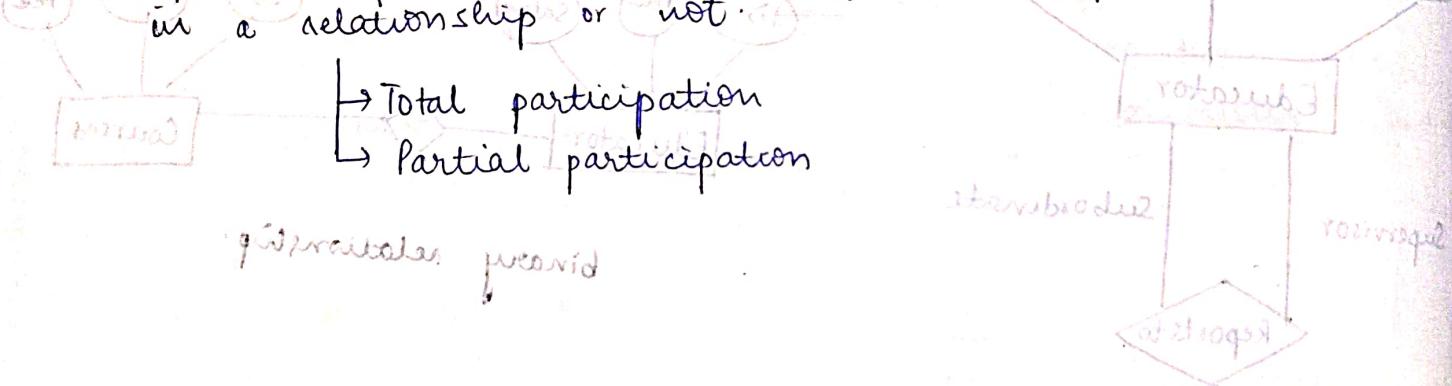
The maximum number of relationships in which an entity can participate is also called mapping cardinality

- One to one relationship
- One to many relationship
- Many to one relationship
- Many to many relationship

Participation constraints

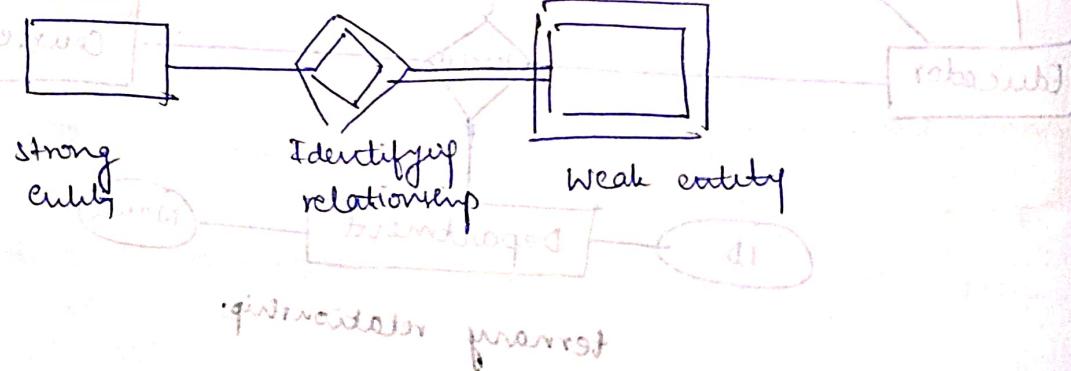
represents if all the members of an entity-set participate in a relationship or not.

- Total participation
- Partial participation



Weak entity - entity which cannot be uniquely identified from its attributes alone.

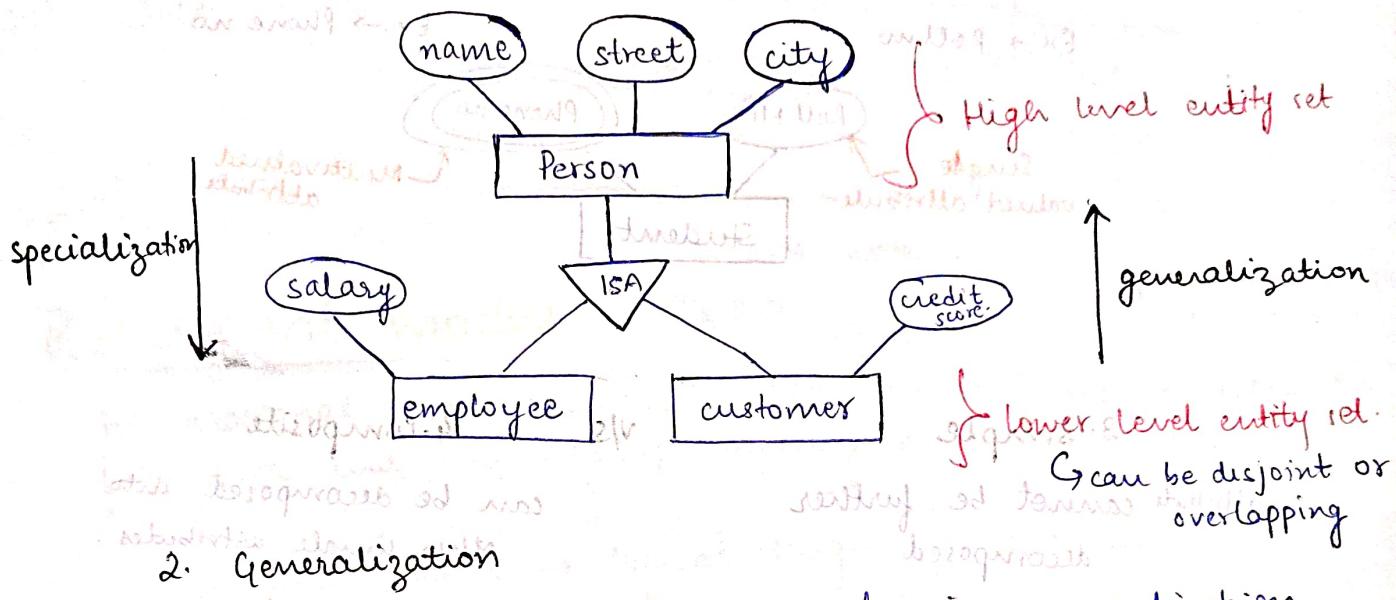
Strong entity - entity which has key attribute.



Extended ER Features

1. Specialization

The process of designating subgroupings within an entity set is called specialization.



2. Generalization

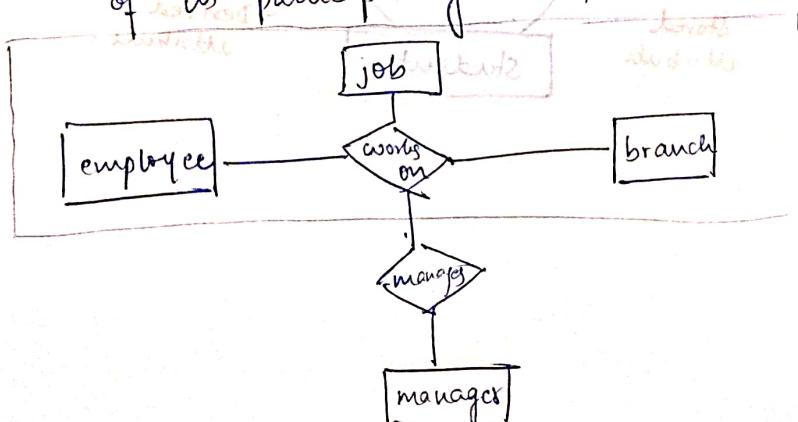
Commonality can be expressed using generalization. Contained relationship b/w higher level entity set & lower level entity set.

Total generalization / specialization - Each higher level entity set must belong to lower level entity set.

Partial generalization / specialization - Some members in higher level entity set may not belong to lower level entity set.

3. Aggregation

Relationship that contains another relationship or one of its participating entity



Types of attributes

1. Single valued

attributes can have only one value.

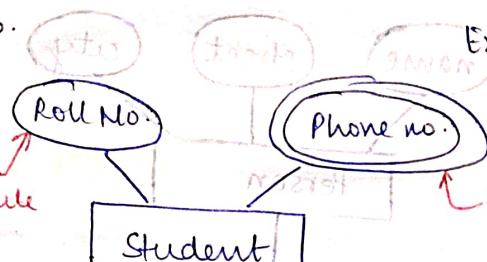
v/s

2. Multi valued

attributes can have more than one value.

Ex → Roll no.

Single valued attribute



Ex → Phone no.

Multivalued attribute

v/s 4. Composite

3. Simple

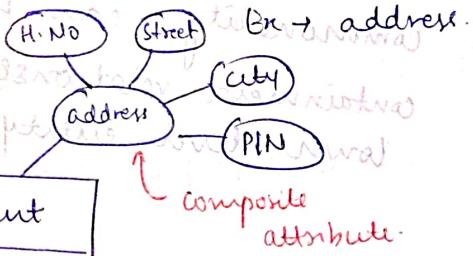
attribute cannot be further decomposed

Ex → age

Simple attribute

Student

can be decomposed into other simple attributes.



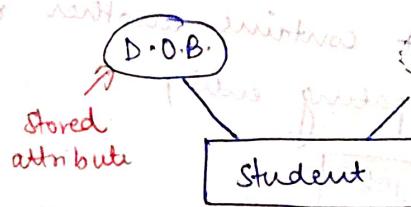
5. Stored

attributes whose value is stored in memory.

6. Derived

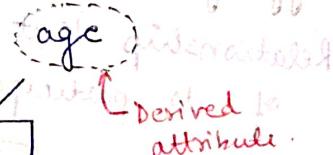
attributes which can be derived from a given attribute.

Ex → D.O.B.



Ex → age

Derived attribute



7. Prime attribute which is part of some key

v/s 8. Non prime attribute attribute which is not part of some key.

9. Complex attributes — Multivalued + Composite

(Ex → address, phone no.)

Relationships → fast processing of relationships of fact statements many words

Relational model

The relational model uses a collection of tables to represent both data and relationships among data.

Database schema → logical design of database.

Relation_name (attribute1, attribute2, ..., attributeN)

UNIQUE
NOT NULL

Domain constraint → specifies the important condition that we want each relation to satisfy.

No two rows in different relations can have same attribute values.

Degree or Arity → number of attributes in a relation.

Cardinality → number of tuples in a relation.

NULL → representation of no value.

Key → attribute/group of attributes that can uniquely identify a tuple in a relation.

Relationships → fast processing of relationships of fact statements many words.

Relationships → fast processing of relationships of fact statements many words.

Relationships → fast processing of relationships of fact statements many words.

Types of keys -

1. Super key - All possible keys of a relation
→ superset of all keys.
2. Candidate key - Minimal superkey i.e. called candidate key.
→ A key whose proper subset is not a key.
3. Primary key - chosen from candidate key for implementation

4. Alternate key -

other than primary key

5. Foreign key -

Key referring to some value in other table
(can be NULL)

Referential Integrity

The values in the foreign key should be present in the column of other table it is referring to

The foreign key should always refer to primary key

Functional Dependencies

Belief $A \rightarrow B$ (B is functionally dependent on A)
 if each value in A is associated with exactly one value
 in B in the relation. Then A also functionally depends on B .

- ① A functional dependency $\alpha \rightarrow \beta$ always holds if α is the key to the relation R .
 $\alpha, \beta \rightarrow$ attributes / set of all attributes in R .

Closure of an attribute -
 set of attributes that can be derived from the given attribute.

Ex $\rightarrow R(A, B, C, D)$

FDS $\rightarrow \{A \rightarrow B, A \rightarrow D, B \rightarrow C\}$

$A^+ = \{A, B, C, D\}$

$B^+ = \{B, C\}$

$C^+ = \{C\}$

$D^+ = \{D\}$

Ex $\rightarrow R(A, B, C, D, E)$

FDS $\rightarrow \{A \rightarrow B, B \rightarrow C, C \rightarrow D, D \rightarrow E, E\}$

$A^+ = \{A, B, C, D, E\}$

$B^+ = \{A, B, C, D, E\}$

$C^+ = \{A, B, C, D, E\}$

$D^+ = \{A, B, C, D, E\}$

$E^+ = \{E, A, B, C, D\}$

First Normal Form

\rightarrow Trivial functional dependencies.

These functional dependencies are satisfied by all relations

Ex $\rightarrow A \rightarrow A$
 $\alpha \rightarrow \beta$ if $\beta \subseteq \alpha$

→ Reflexivity rule

$\rightarrow \alpha \rightarrow \beta$ holds if $\beta \subseteq \alpha$

→ Augmentation rule

if $\alpha \rightarrow \beta$ holds, then $\gamma\alpha \rightarrow \gamma\beta$ also holds

→ Transitivity rule

if $\alpha \rightarrow \beta$ and $\beta \rightarrow \gamma$ holds, then $\alpha \rightarrow \gamma$ also holds.

Armstrong's axioms

Additional Rules

→ Union Rule -

If $\alpha \rightarrow \beta$ and $\alpha \rightarrow \gamma$ holds, then,
 $\alpha \rightarrow \beta\gamma$ also holds.

→ Decomposition Rule -

If $\alpha \rightarrow \beta\gamma$ holds, then, both $\alpha \rightarrow \beta$ and $\alpha \rightarrow \gamma$ also hold.

→ Pseudo transitivity rule.

If $\alpha \rightarrow \beta$ and $\beta \rightarrow \gamma$ holds, then,
 $\alpha \rightarrow \gamma$ also holds.

Closure of a set of functional dependencies.

Finding all possible FDs from given set of FDs

$$\{A, C\} \rightarrow R(A, B, C, D)$$

$$FDs = \{A \rightarrow B, B \rightarrow C, AB \rightarrow D\}$$

$$\{A, C\} \rightarrow R(A, B, C, D)$$

$$\{A, C\} \rightarrow R(A, B, C, D) F^+ = \{A \rightarrow C, A \rightarrow D\}$$

transitive
rule

↪ pseudo transitive rule

$$A \rightarrow A \quad f \quad q \rightarrow q$$

the first fd is
itself

and next fd is
itself

and next fd is
itself

and next fd is
itself

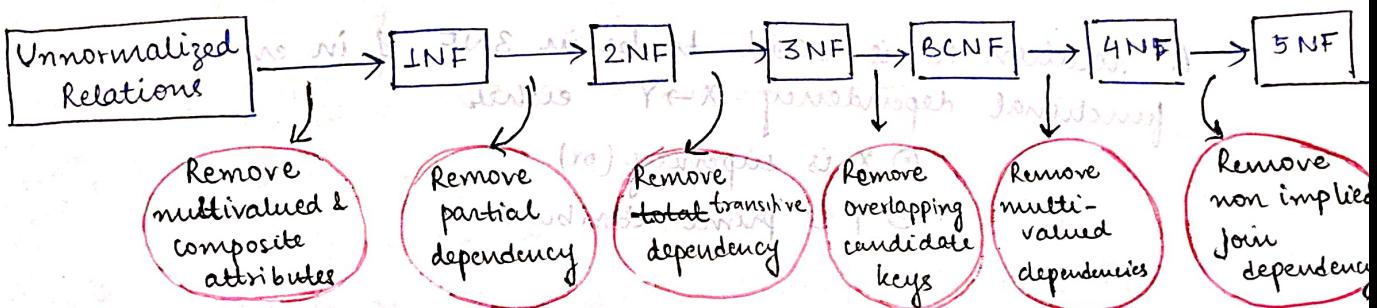
NORMALIZATION

Normalization is the process of minimizing redundancy from a relation or a set of relations.

→ can cause insert, delete, update anomalies.

① Process of grouping attributes into (well) structured relations, that contain minimum redundancy.

Process of normalization



First Normal Form (1NF)

A relation R is said to be in 1NF if there is no any multivalued & composite attribute.

Second Normal Form (2NF)

→ Relation is in 1NF

→ no ^{non} prime attribute in R should be partially dependent on R prime attribute.

Partial dependency :-

Ans $x \rightarrow y$ is partial dependency if some attribute $A \in X - y$ can be removed and the dependency still holds.

→ If candidate key contains more than one attribute, then a part of candidate key should not derive anything else.

primarily good for
minimizing redundancy

* No partial dependency if candidate keys contains only one prime attribute.

Third normal form (3NF)

- Relation R should be in 2NF
- No any non prime attribute in R should be transitively dependent on key of R.

A relation R is said to be in 3NF if in every functional dependency $X \rightarrow Y$ either

- ① X is superkey (or) prime attribute
- ② Y is prime attribute

Boyce Codd Normal Form (BCNF)

- Relation R should be in 3NF
- for every functional dependency $\alpha \rightarrow \beta$, α should be superkey.

* Functional dependency may not be preserved

* Any relation with 2 attributes is always in BCNF

* Normal form considered adequate for normal database design

Ans → 3NF.

→ bcoz dependency preserving lossless-guarantees.

Lossy Decomposition

The decomposition of relation R into $R_1, R_2, R_3, \dots, R_n$ is lossy when the join of R_1, R_2, \dots, R_n does not produce same relation as in R .

< No. of tuples are increased >

If no common attribute is present



$$R = R_1 \bowtie R_2 \bowtie R_3 \dots \bowtie R_n$$

Lossless Decomposition

present then, no common attribute is present for decomposition is lossless.



$$R = R_1 \bowtie R_2 \bowtie R_3 \dots \bowtie R_n$$

Lossy

Decomposition

- * If decomposition is lossless when the common attribute is a key in one of the sub-relations.

* Lossless decomposition is guaranteed upto 3NF.

from
when at least 2 parts
is split into 2 parts
These
algorithm is
used to check
if decomposition
is lossless or
& not lossless
is not lossless

In BCNF, sometimes, we it is possible to satisfy only lossless decomposition property or functional dependency preservation property.

We focus on making decomposition lossless.

∴ Lossless decomposition holds in BCNF also.

Dependency Preserving Decomposition

A decomposition $D = \{R_1, R_2, R_3, \dots, R_n\}$ of R is dependency preserving with respect to a set F of functional dependencies if

$$(F_1 \cup F_2 \cup F_3 \cup \dots \cup F_n)^+ = F^+$$

not ensured
in BCNF

$$(b, 352) \rightarrow 3800$$

∴ we are deriving that losses are in $(b, 352) \rightarrow 3800$.

SQL: Structured Query Language

- ① Domain specific language
- ② used for managing data in RDBMS

→ SQL Data types

CHAR (size) → a fixed length string (can contain letters, numbers, special characters)

size → $0 \leq \text{size} \leq 255$
Default → 1.

VARCHAR (size) → variable length string

size defines the max. length of the string

size → $0 \leq \text{size} < 65535$

size → maximum no. of characters

Bit (size) → a bit value type no. with size

size → number of bits

BOOL → 0 is considered false, non zero value is true

INT (size) → medium integer

BIGINT (size) → large integer

FLOAT (size, d) → floating point no.

d → no. of digits after decimal point.

size → total no. of digits

FLOAT(p) → floating point no.

p is used to determine whether to use FLOAT or DOUBLE

$0 < p < 24 \rightarrow \text{float}$

$25 < p < 53 \rightarrow \text{double}$

DOUBLE (size, d)

DECIMAL (size, d) → an exact fixed point no.

* SQL is not case sensitive

1. SELECT COMMAND

used to retrieve data from one or more tables.

SELECT <column-name> FROM <table-name>;

* returns
duplicates
if present.

(most tables have NOT NULL constraint)

SELECT * FROM <table-name> → returns whole table.

SELECT DISTINCT <field> FROM <table>;

(will return unique records)
always used only with SELECT

SELECT DISTINCT <field> FROM <table>

→ distinct combination
of the 2 fields

2. WHERE Clause

① used with SELECT, UPDATE, DELETE & INSERT commands.

② used to filter specific rows from the table

SELECT <column name> FROM <table-name>

WHERE <condition>

③ Relational operators that can be used in WHERE clause -

=, <, >, <=, >=

④ Logical operators -

AND, OR, NOT

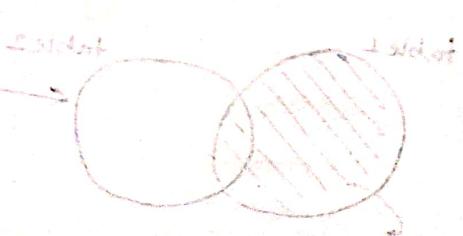
⑤ BETWEEN operator -

Used to filter the records in specific range.

Range - <start> BETWEEN LB and UB (both inclusive)

start must be greater than or equal to LB

range numbers



SQL Joins

- ① needed to retrieve records from more than one table collectively
- ② 2 tables must have one common column.

Types of joins —

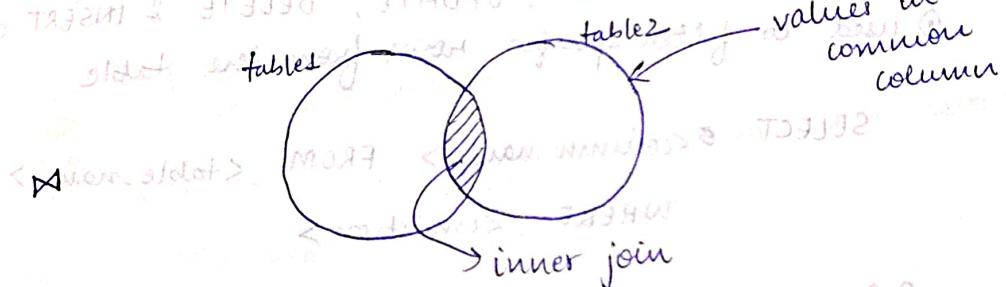
- Full Join (Full outer join)
- Left Join (left outer join)
- Right Join (right outer join)
- Inner Join (Equi-Join) or simple join
- Self Join

No. of attributes in result = $\text{Deg}(\text{table1}) + \text{Deg}(\text{table2})$

1. Inner join

$\text{SELECT * FROM } <\text{table1}> \text{ INNER JOIN } <\text{table2}> \text{ ON }$

$<\text{table1}>. <\text{column}> = <\text{table2}>. <\text{column}>$



equivalent to

$\text{SELECT * FROM } <\text{table1}>, <\text{table2}> \text{ WHERE }$

$<\text{table1}>. <\text{column}> = <\text{table2}>. <\text{column}>$

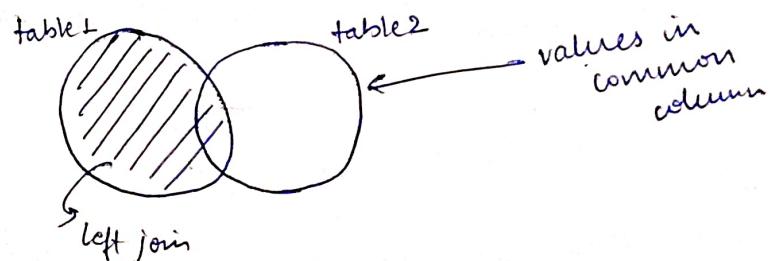
2. Left Join

$\text{SELECT * FROM } <\text{table1}> \text{ LEFT JOIN } <\text{table2}> \text{ ON }$

$<\text{table1}>. <\text{column}> = <\text{table2}>. <\text{column}>$



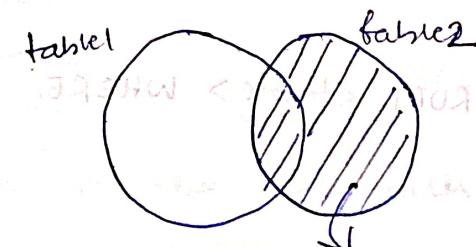
common tuples + all tuples from table2



3. Right join

SELECT * FROM <table1> RIGHT JOIN
 <table2> ON <table1>. <column> = <table2>. <column>

Shows all the tuples from table 2 + common tuples from both tables.

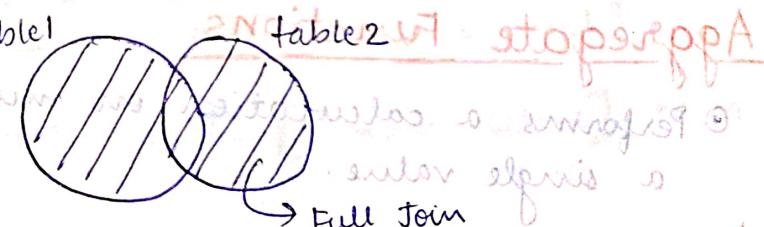


Right join = T3E1J2 *

4. Full join

SELECT * FROM <table1> FULL JOIN <table2>
 ON <table1>. <column> = <table2>. <column>

SELECT common tuples + tuples from table1 + tuples from table2



Full join = T3E1J2

5. Self join

only 1 table is involved in this join

self join is not keyword

SELECT * FROM <table1> AS <t1>, <table1> AS <t2>
 WHERE <t1>. <column> = <t2>. <column>

(t1) and (t2) are temporary table names.

t1	t2	t1	t2
1	1	1	1
1	2	1	2
2	1	2	1

~~most useful command standard is MySQL & PostgreSQL~~

~~SQL SELECT clause <table> WHERE <condition>~~

LIMIT clause

used to limit the number of fetched records from a huge database table.

SELECT * FROM <table> WHERE <condition> LIMIT n

~~imp~~

* SELECT * FROM <table> WHERE <condition> LIMIT m, n

<table> ~~MOST~~ max n records after leaving first m records

<table> <set> = leaving first m records

most output from a single table

leave first

MOST +

return n records

Aggregate Functions

① Performs a calculation on multiple values and returns a single value.

* ② can only be used with SELECT clause

- MIN(<column>) — returns min value in column
- MAX(<column>) — returns max value in column
- SUM(<column>) — returns sum of values in column
- COUNT(<column>) → returns no. of records ~~not~~
- AVG(<column>) → returns average of values in column

all aggregate functions ignore NULL values except count(*)

A	B
1	N
2	N
N	N
3	N

$$\text{count}(*) = 4$$

$$\text{count}(A) = 3$$

$$\text{count}(B) = 0$$

$$\text{SUM}(A) = 6$$

$$\text{AVG}(A) = 2$$

$$\text{MAX}(A) = 3$$

$$\text{MIN}(A) = 1$$

$$\text{SUM}(B) = \text{NULL}$$

$$\text{AVG}(B) = \text{NULL}$$

$$\text{MAX}(B) = \text{NULL}$$

$$\text{MIN}(B) = \text{NULL}$$

IN operator

- ① The IN operator allows you to specify multiple values in a WHERE clause.
- ② IN operator is shorthand for ~~not~~ multiple OR conditions.

~~multiple OR conditions~~ → ~~multiple OR conditions~~ → ~~multiple OR conditions~~

SELECT * FROM <table> WHERE <column> IN (value1, value2...)

- ③ IN operators is also used with subqueries.

SELECT * FROM <table> WHERE <column> IN (SELECT <column> FROM <subquery> WHERE ...)

ALIAS

- ① used to give temporary name to a table, or a column of a table.

- ② only exists during duration of that query.

- ③ created using AS.

[space is not allowed in column name of a table.
To print space, ALIAS can be used]

SELECT CustomerName AS [Customer Name] FROM Customers.

GROUP BY Clause

- ① used in collaboration with SELECT statement to arrange identical data into groups.

- ② all attributes used in GROUP BY clause must be present in the SELECT clause.

- ③ attributes not present in GROUP BY clause can appear in aggregate functions.

SELECT DNo, AVG(Salary) FROM Employees

GROUP BY DNo;

→ (Salary of D1) → (Salary of D2) → (Salary of D3) → (Salary of D4)

→ (Salary of D1) → (Salary of D2) → (Salary of D3) → (Salary of D4)

→ (Salary of D1) → (Salary of D2) → (Salary of D3) → (Salary of D4)

→ (Salary of D1) → (Salary of D2) → (Salary of D3) → (Salary of D4)

HAVING clause

similar to WHERE clause.
can only be used with GROUP BY clauses.

SELECT DNO, AVG(Salary) FROM Employees

GROUP BY DNO HAVING AVG(Salary) > 50.

Order of execution → WHERE → GROUP BY → HAVING

WHERE checks the condition on individual rows

Having checks the condition on a group of rows.

Having cannot be used with aggregate functions

can be used with aggregate functions

executed before GROUP BY clause

executed after GROUP BY clause

Set operators

Set Operators

Set operators are used to combine results of two or more SELECT statements.

no. of columns returned from the

individual queries must be equal.

(SELECT * FROM Table1) <set operator> (SELECT * FROM Table2)

- UNION (duplicate rows are eliminated)
- UNION ALL (duplicate rows are retained)
- INTERSECT (selects only common rows)
- MINUS / EXCEPT (set difference)

ORDER BY

- ① ORDER BY clause is used to sort the result set in ascending or descending order.

priyanshu was algorithm in how protocal
ASC/DESC default

```
SELECT * FROM <table-name> ORDER BY <column-name> ASC/DESC
```

- ② default order of sorting → ascending order

- ③ ORDER BY clause is executed in the very end.

(priyanshu T37302 Product NO97 * T37372)

(Singapore most popular T37372)

Subquery

- ① also called inner query, nested query

- ② query within a SQL query that is embedded within the WHERE clause

2 types of subquery

→ Single Row subquery

→ Multiple Row subquery.

Single Row subquery

- * Product with second highest price

SELECT MAX(price) FROM Products

WHERE price < (SELECT max(price) FROM Products)

- * Product with third highest price

SELECT MAX(price) FROM Products

WHERE price < (SELECT max(price) FROM products)

WHERE price < (SELECT MAX(price) FROM products));

Relational operators can be

used in single row subqueries.

Multiple Row subquery

Operators used in Multiple Row subquery

1. IN operator

- Allows to specify multiple values in WHERE clause.
- Shorthand for logical OR.

SELECT * FROM Customers WHERE country IN
 (SELECT country FROM suppliers);

* Find all customers who have placed more than 2 orders -
 2 tables
 Customers, Orders
 SELECT * FROM Customers WHERE customerID IN
 (SELECT customerID FROM Orders
 GROUP BY customerID HAVING count(customerID) > 2)
 pregrouped by customerID
 pregrouped by customerID

2. ANY operator

SELECT <column-name(s)>
 FROM <table-name>
 WHERE <column-name> operator ANY

standard comparison operator
 (= <> <= > < >)

standard comparison operator

3. ALL operator

SELECT <column-name(s)>
 FROM <table-name>
 WHERE <column-name> operator ALL
 (subquery)

4. EXISTS operator

- ① used to test for the existence of any record in a subquery.
- ② returns true if subquery returns one or more records.

```
SELECT <column_name(s)>
FROM <table_name>
WHERE <column_name> EXISTS
      (subquery)
```

* 'SELECT NULL' returns 1 record with column & row NULL

null
null

• did not go into the relations given below at all

(R) <(2) ROLL, ROLLNO>

RELATIONAL ALGEBRA

- Relational algebra is a procedural query language which takes instances of relations as input & yields instances of relations as output.
- provides theoretical foundation for relational databases

Basic operators—

- | | |
|-----------------------|----------------------------------|
| ① Select (σ) | ② Intersection (\cap) |
| ③ Project (π) | ④ Cartesian Product (\times) |
| ⑤ Union (\cup) | ⑥ Rename (ρ) |
| ⑦ Set Difference (-) | ⑧ Division (\div) |

1. Selection operator

- fetches tuples which satisfy a condition
- $\pi_{\langle \text{conditions} \rangle} (\langle \text{Relation-name} \rangle)$
- similar to WHERE clause of SQL

① Selection operator is commutative

$$\pi_{\langle \text{cond}_2 \rangle} (\pi_{\langle \text{cond}_1 \rangle} (R)) = \pi_{\langle \text{cond}_1 \rangle} (\pi_{\langle \text{cond}_2 \rangle} (R))$$
$$= \pi_{\langle \text{cond}_1 \rangle \wedge \langle \text{cond}_2 \rangle} (R)$$

2. Projection operator

- used to select only certain columns of the table

$$\pi_{\langle \text{column-name(s)} \rangle} (R)$$

* eliminates duplicates in result.

- similar to SELECT statement `SELECT DISTINCT`

RELATIONAL ALGEBRA

most relations of several groups of operations are developed based on set operations

3. Set operations

- ① the 2 relations on which set operations are applied should necessarily have similar data types of the tuples of relations

- ② they should be union compatible

(i) ~~not present~~ are not present.

(x) domain restrictions

i) Union operator (\cup)

commutative & associative

ii) Intersection operator (\cap)

iii) Set Difference (-)

→ difference of two sets is the set of elements which are in the first set but not in the second set

→ if $R \cap S = \emptyset$, then $R - S = R$ and $S - R = \emptyset$

→ if $R \cap S \neq \emptyset$, then $R - S \neq R$ and $S - R \neq \emptyset$

→ if $R \cap S = R$, then $R - S = \emptyset$ and $S - R = S$

- * $\pi_x(R_1 \cup R_2) = \pi_x(R_1) \cup \pi_x(R_2)$
- * $\pi_x(R_1 \cap R_2) \neq \pi_x(R_1) \cap \pi_x(R_2)$

4. Cartesian Product

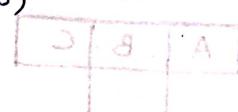
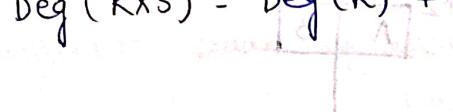
cross product

$$|R \times S| = m \cdot n$$

$$\text{Deg}(R \times S) = \text{Deg}(R) + \text{Deg}(S)$$

$$|R| = m$$

$$|S| = n$$



→ Joins

are nothing but joins

i) Conditional join

$R \bowtie S$ for various joins between $R_1 \bowtie_{\text{condition}} R_2$

ii) Equi join

$R_1 \bowtie_{R.c = S.c} R_2$

iii) Natural join

Both relations should have atleast one common column name

$R \bowtie S$

iv) left outer join

$R \bowtie L S$

v) Right outer join

$R \bowtie R S$

vi) Full outer join

$R \bowtie S$

5. Rename operator

① similar to aliasing in SQL.

② used to give temporary name to relation / attribute

$\rho_{\text{newname}}(\langle \text{relation} \rangle) \rightarrow \text{Rename relation}$

$\rho_{R(A,B,C)}(R) \rightarrow \text{Renames columns to } A, B, C$

Q

$$(a^2) \times \pi \cap (1^2) \times \pi = (a^2 \times 1^2) \times \pi =$$

6. Division operator

$$\textcircled{O} R_1 \div R_2$$

\textcircled{O} Attribute set of $R_2 \subseteq$ Attribute set of R_1 .
Attribute set of R1 is greater than or equal to attribute set of R2.

$$\text{Result set} = (\text{Attribute set of } R_1) - (\text{Attribute set of } R_2)$$

$$R_1 \quad \begin{array}{|c|c|c|} \hline A & B & C \\ \hline \end{array} \quad R_2 \quad \begin{array}{|c|c|} \hline A & B \\ \hline \end{array} = \text{Def}(R_1 \setminus R_2)$$

$R_1 \div R_2 \rightarrow$ Those values of C which are associated with every value of A, B in R_2

<used in those queries which contain 'for all' >

$\exists x \in R$

exist value of x such that all values of y in R

$\exists x \in R$

exist value of x such that all values of y in R

$\exists x \in R$

exist value of x such that all values of y in R

$\exists x \in R$

exist value of x such that all values of y in R

$\exists x \in R$

exist value of x such that all values of y in R

$\exists x \in R$

exist value of x such that all values of y in R

$\exists x \in R$

exist value of x such that all values of y in R

$\exists x \in R$

exist value of x such that all values of y in R

$\exists x \in R$

exist value of x such that all values of y in R

TRANSACTION

logic unit of database which includes one or more database operations

- Commit - command used to make the changes permanent in a database. (Operation successful)
 - Rollback - command used to revert back to the original state of database (Operation unsuccessful)
- Database state becomes same as it was on last commit

ACID Property of Transaction

1. Atomicity :- All or none. (either complete transaction or never start it)
2. Consistency :- the database state must be consistent
3. Isolation :- transactions should execute such that it does not affect nor is affected by any other concurrent transaction
4. Durability :- changes made should be reflected back into database (permanently)

Schedule-

Execution of multiple transactions on a same database concurrently is called schedule.

Why concurrency?

TRANSACTIONS

1. Improved throughput (No. of lines executed per unit time)
2. Resource utilization
3. Reduced waiting time

Problems with concurrency -

1. Recoverability problems

2. Deadlock of busy waiters & - deadlock

3. Serializability issues (state consistency)

Serializability problems ~~convert to utopian ACID~~

→ Lost update problem

→ Dirty read problem ~~optimistic~~

(脏数据问题在并发读写时可能会发生)

→ Phantom tube

→ Unrepeatable read problem

(不可重复读问题)

→ Incorrect summary problem.

Li test über etwars blauje und das wort ~~mittocht~~ - 2

rechts fwo für betzoffe in ram doppf tez wob

Good schedule - concurrent schedules which provide consistent

result after all threads have terminated.

Bad schedule - These schedules do not give consistent

result (suffer from one or more of the above problems)

Arbeit mit den Konkurrenzschichten ~~abro~~ - schubert

• ältere Schichten mit höherer Priorität

• jüngere Schichten sind niedriger priorisiert

• Schichten können sich gegenseitig blockieren

Serial schedules

- ① schedules which do not have interleaving
- ② always give consistent result
- ③ For n transactions, no. of serial schedules = $n!$

Non serial schedules

- ① schedules which have interleaving are called non serial schedule.

edge of interleaving

$T_1 \rightarrow n_1$ steps

$T_2 \rightarrow n_2$ steps

$T_n \rightarrow n_n$ steps

NO. of schedules possible

$$= \frac{(n_1 + n_2 + n_3 + \dots + n_n)!}{n_1! n_2! n_3! \dots n_n!}$$

Serializable schedule

- ① A concurrent schedule which provides result as a serial schedule is called serializable schedule.

Serializability

→ conflict serializability

→ View serializability

1. Conflict Serializability

→ 2 DB statements have conflict if:

→ Both are in different transactions

→ Both are accessing same data item

→ One of them is write operation

→ 2 transactions are conflict equivalent if the conflicting statements appear in the same order in both of them.

→ If one of the 2 transactions is serial schedule, the other one is called conflict serializable schedule.

To prove if a given schedule is conflict serializable,
we use **precedence graph**

- vertices → transactions
- edges → conflicts
- directed graph

T_1	T_2	T_3
$R(x)$		
	$W(x)$	
		$R(y)$
	$W(y)$	
$R(y)$		
		$R(y)$



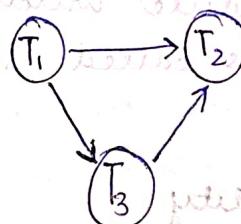
conflict serializable

∴ cycle is present in precedence graph

∴ given schedule is not conflict serializable.

eg. $T_1 \rightarrow T_2 \rightarrow T_3 \rightarrow T_1$

T_1	T_2	T_3
$R(x)$		
	$R(y)$	
		$R(y)$
	$W(x)$	
		$W(x)$
	$R(x)$	
	$R(x)$	



conflict serializable

no cycle

∴ conflict serializable schedule

T_1, T_3, T_2

T_1	T_2	T_3	T_4
$R(A)$			
$W(B)$			
	$R(B)$		
		$R(C)$	
			$R(A)$
			$W(A)$
$W(A)$			
		$R(C)$	
			$W(B)$
			$R(B)$
			$W(C)$



conflict serializable

conflict serializable schedule

conflict serializable

conflict serializable

T_1, T_3, T_4, T_2

T_1, T_3, T_2, T_4

T_1, T_2, T_3, T_4

T_2, T_1, T_3, T_4

2. View Serializability

3 points to remember

3 points
who reads from database

↳ who reads from other's written value

→ who writes last

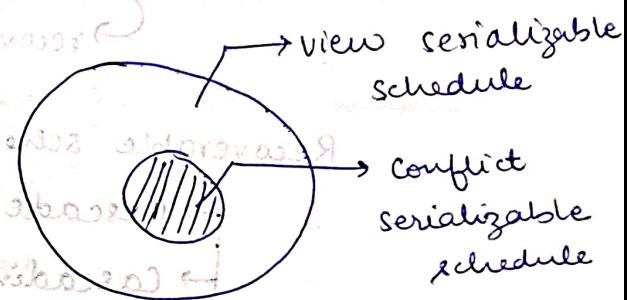
② schedules s_1 & s_2 are view equivalent if both are following similar sequence/rules of view serializability.

① All conflict serializable

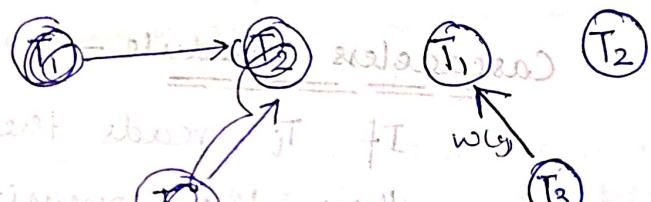
Schedules are view

~~deadlock~~ serializable.

absolutly shocked at how much I have learned



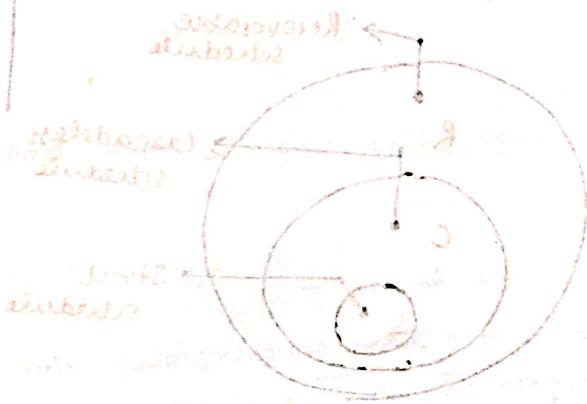
T_1	T_2	T_3
$R(X)$	$R(Y)$	$w(Y)$
$w(Y)$	T	$w(X)$
T	$w(X)$	$w(Y)$



not view serializing

$$\begin{array}{l} T_3 \\ \hline \end{array}$$

3 possible serial schedules



Recoverability

① When no any committed transaction needs to roll back, it is called recoverable schedule.

② If a transaction T_i reads dirty value written by T_j and then the commit operation of T_i should be after the transaction T_j commits.

↳ Recoverable schedule

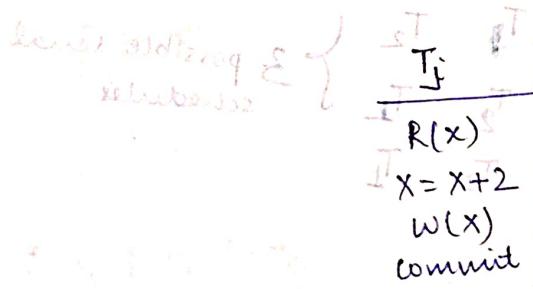
↳ Recoverable schedule using no rebuffers

↳ Cascadeless Recoverable schedule Rollback

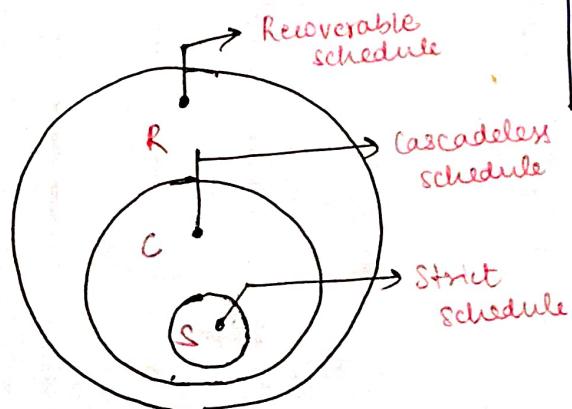
↳ Cascading Recoverable schedule Rollback

③ Cascadeless schedule -

If T_i reads the data item written by T_j then, the commit operation of T_j should appear before read operation of T_i .



$R(x)$
 $x = x+3$
 $w(x)$
commit



Obtaining serializability (view serializable & conflict serializable) is also not practically possible. ∴ locking protocols are used.

Locking Protocols

→ Shared lock (other transactions can share)

~~Lock~~ → shared with other processes - taken for read by other processes

Exclusive lock (no other transaction can read/write)

		Demand currently held ↓	Shared	Exclusive	
		Shared	Shared	X	(X) 2-20
		Exclusive	X	X	(X) 9 (X) 7-2001
ST	IT				
(X) 9	to		(X) 9		(X) 9
(X) 7					(X) 7-2001

Table maintained by DBMS

	Availability	Mode	Count	Blocked
X	0 if available + if lock is held.	S → shared EX → Exclusive	→ 49.2 added when borrowed → also always triggered → No. of transactions holding lock (only in case of shared lock)	→ 49.2 added when borrowed

transactions (FIFO order) that are waiting for the data items to be unlocked.

→ Simple locking mechanism can lead to unrepeatable read problem

T_1	T_2
$\text{Lock_S}(x)$	$\text{LOCK_X}(x)$
$R(x)$	$W(x)$
$\text{Unlock}(x)$	$\text{Unlock}(x)$

To solve unrepeatable read problem, several protocols / ~~rules~~ are developed.

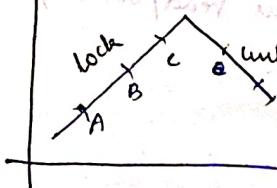
2 Phase Locking Protocols (Basic 2PL)

- Systematic locking mechanism.
- Once unlock is done, a transaction is not allowed to lock any other data item.

Each transaction has 2 phases -

→ Growing phase - acquiring of locks

→ Shrinking phase - releasing of locks.



Serializability

Second

Interference

LOCK-S(X)

R(X)

LOCK-EX(Y)

Unlock(X)

W(Y)

Unlock(Y)

T₂

LOCK-S(X)

R(X)

LOCK-EX(Y)

W(X)

Unlock(X)

Unlock(Y)

Only 2 possible schedules
there are three -

T₁ & T₂ X & Y

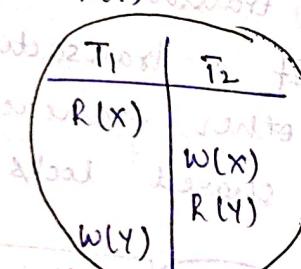
or

T₁ | T₂

R(X)
W(Y)

R(Y)
W(X)

W(X)
W(Y)



schedule
not
allowed
by basic
2PL.

Every schedule which is
allowed under basic 2PL is
conflict serializable also.

Question asked - Tell whether the given
schedule is allowed under 2PL
or not.

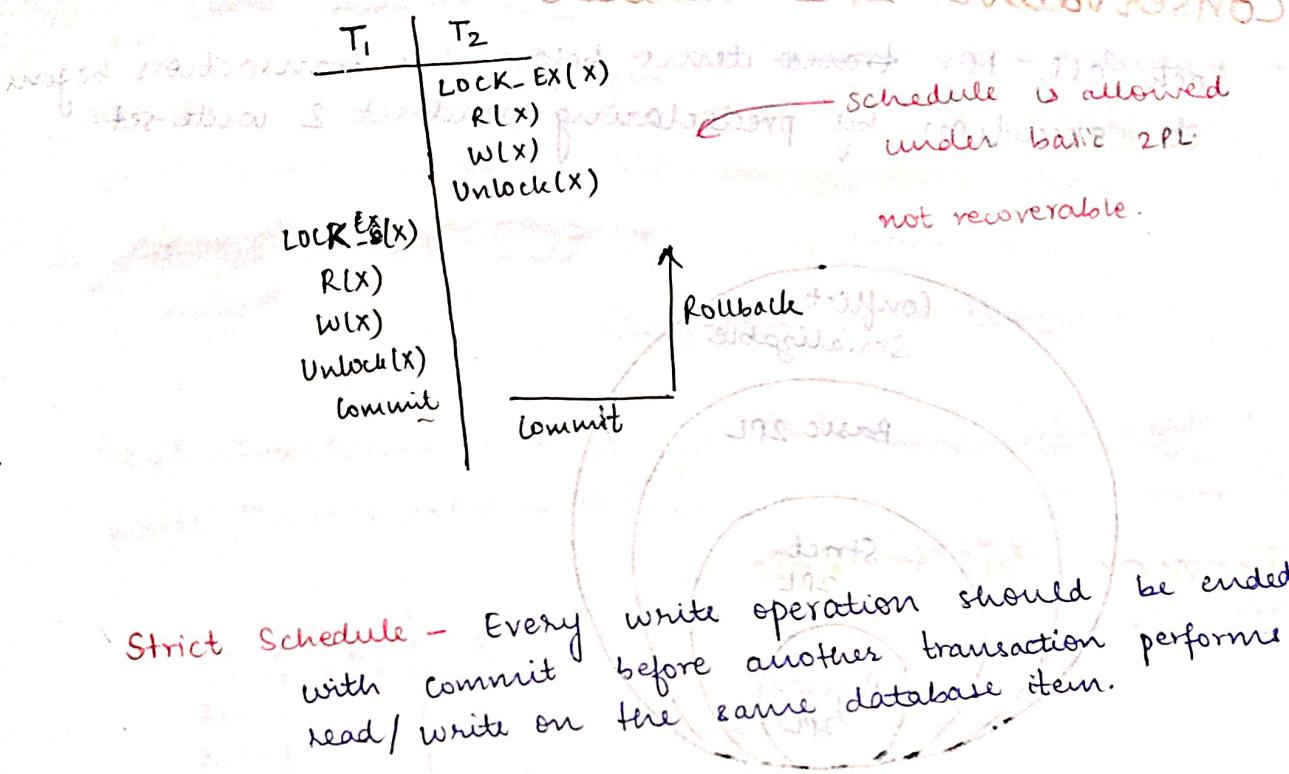
T ₁	T ₂	T ₃	T ₄
W(X)			
W(Y)			
W(Z)			
	R(X)		
	R(Y)		
	R(Z)		

* 2PL suffers from starvation.

* 2PL suffers from deadlock.

T ₁	T ₂
LOCK-S(X)	
R(X)	
	LOCK-S(Y)
	R(Y)
	LOCK-EX(Y)
	W(Y)
	LOCK-EX(X)
	W(X)

* 2 PL can allow non recoverable schedules



Strict Schedule - Every write operation should be ended with commit before another transaction performs read/write on the same database item.

Strict 2PL Protocol

Exclusive locks should be released only after commit.

* Strict 2PL allows only strict schedules.

Rigorous 2PL Protocol

Every lock should be released after commit.

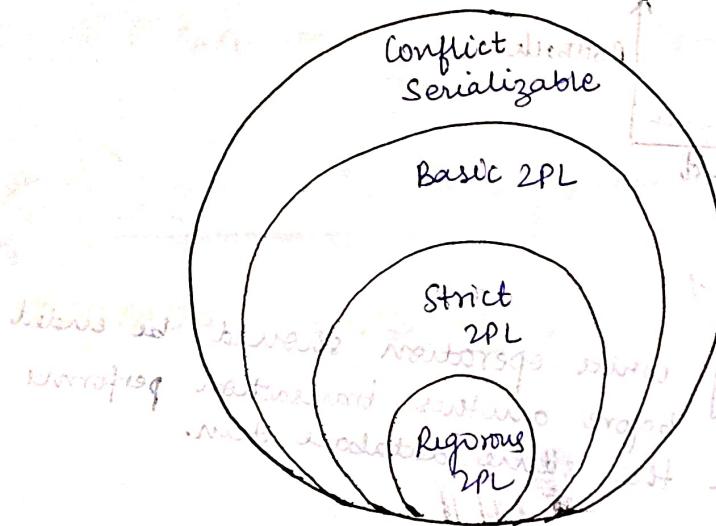
① Strict & Rigorous 2PL do not suffer from dirty read problem.

② Basic 2PL may allow non recoverable schedule.

③ Strict & Rigorous 2PL have only recoverable schedules (cascadeless)

Conservative 2PL - Protocol

Lock all the trans items before the transaction begins its execution by predeclaring (read-set & write-set)



Final notes: five levels of 2PL protocols
differences between them

five levels of 2PL protocols
difference between them

five levels of 2PL protocols
difference between them

best part: may refuse to do 2PL managed 2PL
transaction

student answer: don't allow new writes from 2PL and 3PL

student answer: only 2PL managed 2PL

(interprocess) interleaving

Timestamp Based Protocols

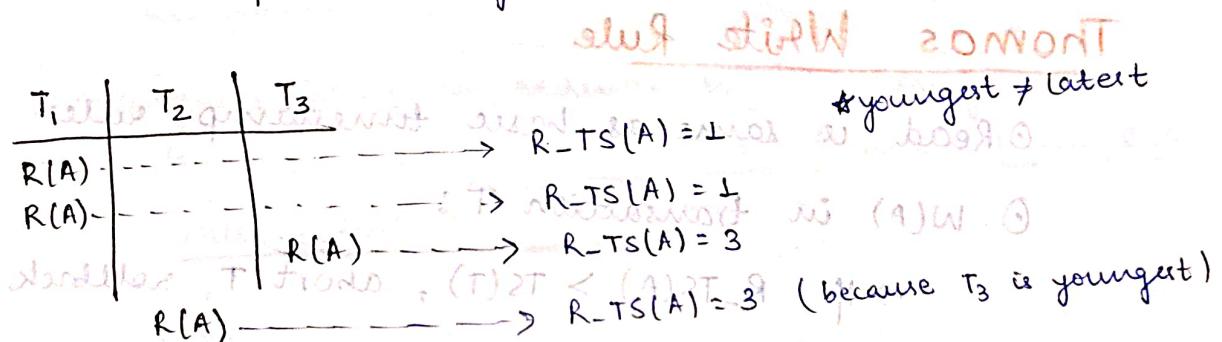
Timestamp:- unique identifier created by DBMS to identify the relative starting time of a transaction.

$TS(T_1) < TS(T_2) \rightarrow T_1 \text{ starts before } T_2$

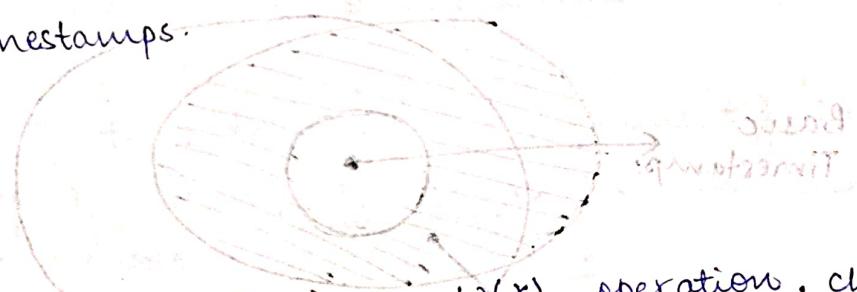
$(T_1 \text{ is older than } T_2)$

Read Timestamp(A) \rightarrow Youngest transaction who reads A

Write Timestamp(A) \rightarrow Youngest transaction who writes A



- (T)2T < (A)2T_W
- Basic timestamp algorithm is used to control concurrency by allowing transactions to run in their relative ordering of timestamp.



When a transaction T issues $w(x)$ operation, check the following conditions:

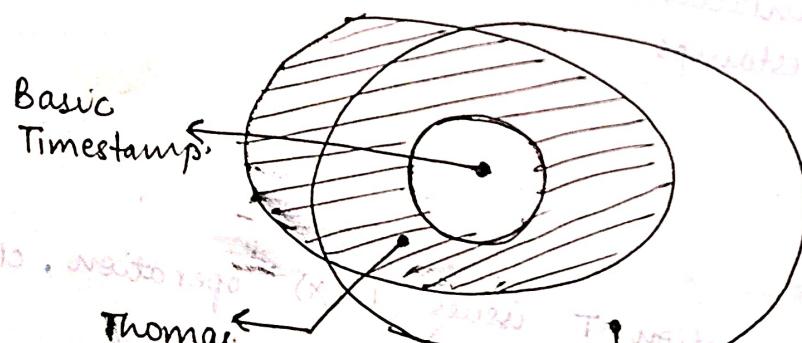
- If $R_TS(x) > TS(T)$ or $W_TS(x) > TS(T)$ abort the transaction and rollback T.

- else, execute $w(x)$ and update $W_TS(x)$ to $TS(T)$.

- When transaction T issues $R(x)$ operation, not check the following conditions -
- ① If $W_{TS}(x) > TS(T)$ then abort and reject the operation.
 - ② else execute $\theta R(x)$ and set $R_{TS}(x)$ to $\max(TS(T), R_{TS}(x))$

Thomas Write Rule

- ① Read is same as basic timestamp rules.
- ② $W(A)$ in transaction T :
 - \rightarrow If $R_{TS}(A) > TS(T)$, abort T , rollback
 - \rightarrow else if $W_{TS}(A) > TS(T)$, skip the write operation
 - \rightarrow else perform $W(A)$ and update $W_{TS}(A) = TS(T)$

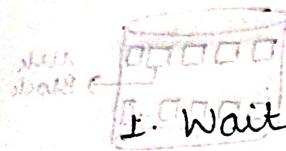


- ③ Some schedules which follow Thomas write rule are not conflict serializable

Deadlock Prevention

1. Wait-DT

Let T_i = older transaction. T_j = newer younger transaction
 $TS(T_i) < TS(T_j)$



1. Wait Die -

→ If T_i tries to acquire lock currently held by T_j , then, T_i waits till the lock is available.

→ If T_j tries to acquire lock currently held by T_i , then, T_j is aborted and restarted with same time stamp.

2. Wait Wound -

→ If T_i tries to acquire lock currently held by T_j , then, T_j is pre-empted by aborting it.

→ If T_j tries to acquire lock currently held by T_i , then, T_j waits till the lock is available.

Resource	Status	Owner	TS
R1	Available		
R2	Locked	T1	TS(T1)

SM & S = browser for registration

SM & S = browser for message transmission

SM & S = browser for message reception

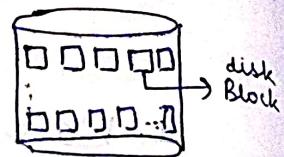
(browser for session initiation, termination, modification, cancellation)

SM & S = browser for session initiation, termination, modification, cancellation

Disk Block - logical unit of storage created by file system.

- ① logical unit of storage created by file system.
- ② A combination of 1-2 sectors.
- ③ size = 512B to 2KB





blue patterns that emerged at night it flew

Record storage - didn't went . IT pd
Sister hit the

Suppose, 1 disk block = 64 bytes = 2^6 bytes

Name	Roll No.	...

$$\Rightarrow 1 \text{ record} = 4 \text{ bytes} \quad \text{No. of records in 1 disk block} = \frac{2^6}{2^2} = 2^{4-2} = 16$$

If the table contains 512 records, this is

total number of disk blocks needed = $\frac{512}{16} = \underline{\underline{32 \text{ blocks}}}$

privado per befreiung si T. west. T. per

Fixed v/s Variable Length Records -

Roll No.	Name	Photo (max. 2 MB)	Signature (max. 200 KB)

max. size of 1 record \approx 2.2 MB

If a student uploads pic of size 1 MB, size 200 kb,
size of record = 1.2 MB.

In fixed length record, size of record is fixed = 2.2MB
(extra space is wasted)

In variable length record, size of record is not fixed
(space ~~savd~~ utilized)

Spanned Mapping

record of a file is stored ~~in~~ inside the block even if it can only be stored partially, hence, record is spanned over multiple blocks.



Adv :- no wastage of memory
(internal fragmentation)



Disadv :- memory access time is larger

Block size = 32B

Record size = 5B

6 records in block 1

2B of 7th record in b1

3B of 7th record in b2

Unspanned Mapping

record of file is stored inside a block only if it can be completely stored inside it.



Adv :- access time of records is less



Disadv :- wastage of memory
(internal fragmentation)

Block size = 32B

Record size = 5B

6 records in B1

7th records in B2

2B in B1 wasted

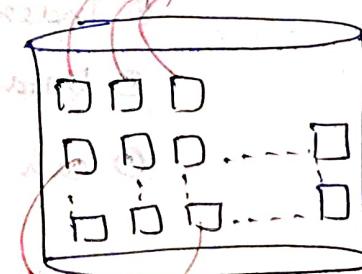
Indexing

index file is created to store information of which records are present in which disk block.

Index file / table is also stored in disk blocks.

→ Dense index:

Index record is for each database record.



→ Sparse Index:

Index record will be for a few few records only in next pages.

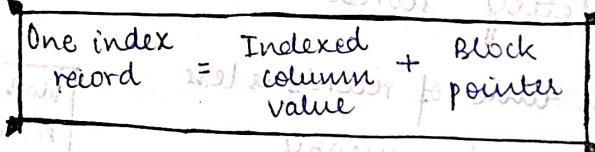
DB records.

Indexing Techniques

- Primary indexing
- clustering indexing
- Secondary key indexing
- Secondary non key indexing

1. Primary indexing

- ① Indexing is done on key primary key or super key.
- ② Data must be ordered on index.
- ③ It's always sparse index



2. Clustering indexing

- ① Indexing is done on non key field. (duplicate values may be present)
- ② Data must be ordered on index field

→ Indexing is done for unique value of non key

③ sparse indexing

3. Secondary key indexing

- ① Indexing done on primary key or any super key
- ② Data must not be ordered on index field
- ③ can be dense or sparse

<non key> → similar to key indexing
except that indexing is done on non key field.

B Tree

- ① Tree based indexing technique.
- ② Dynamic indexing technique.
- ③ Self balancing tree search tree

→ grows horizontally as well as vertically.

Order of B tree - max children of a node of B Tree

(order is denoted by 'p')

An order p B-Tree -

↳ every node has max p children (tree pointers)

↳ every node other than root should have atleast $\lceil \frac{p}{2} - 1 \rceil$ nodes.

↳ every node can have atmost $(p-1)$ keys and p tree pointers

↳ Root can have minimum 1 node key

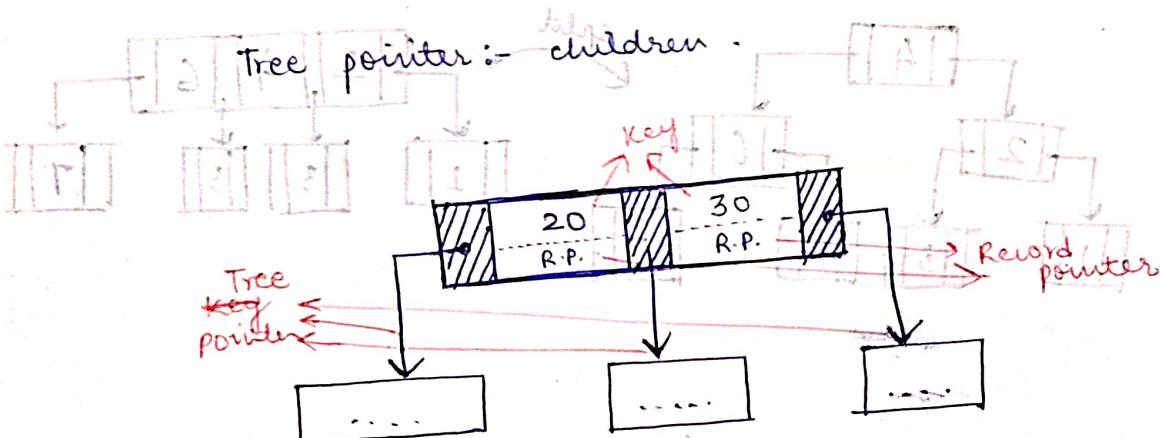
↳ all leaves are at same level.

B Tree node structure

Key - index pointer

Record pointer:- pointer to the disc block where key index is stored.

Tree pointer :- children



Example -

order 3 B Tree

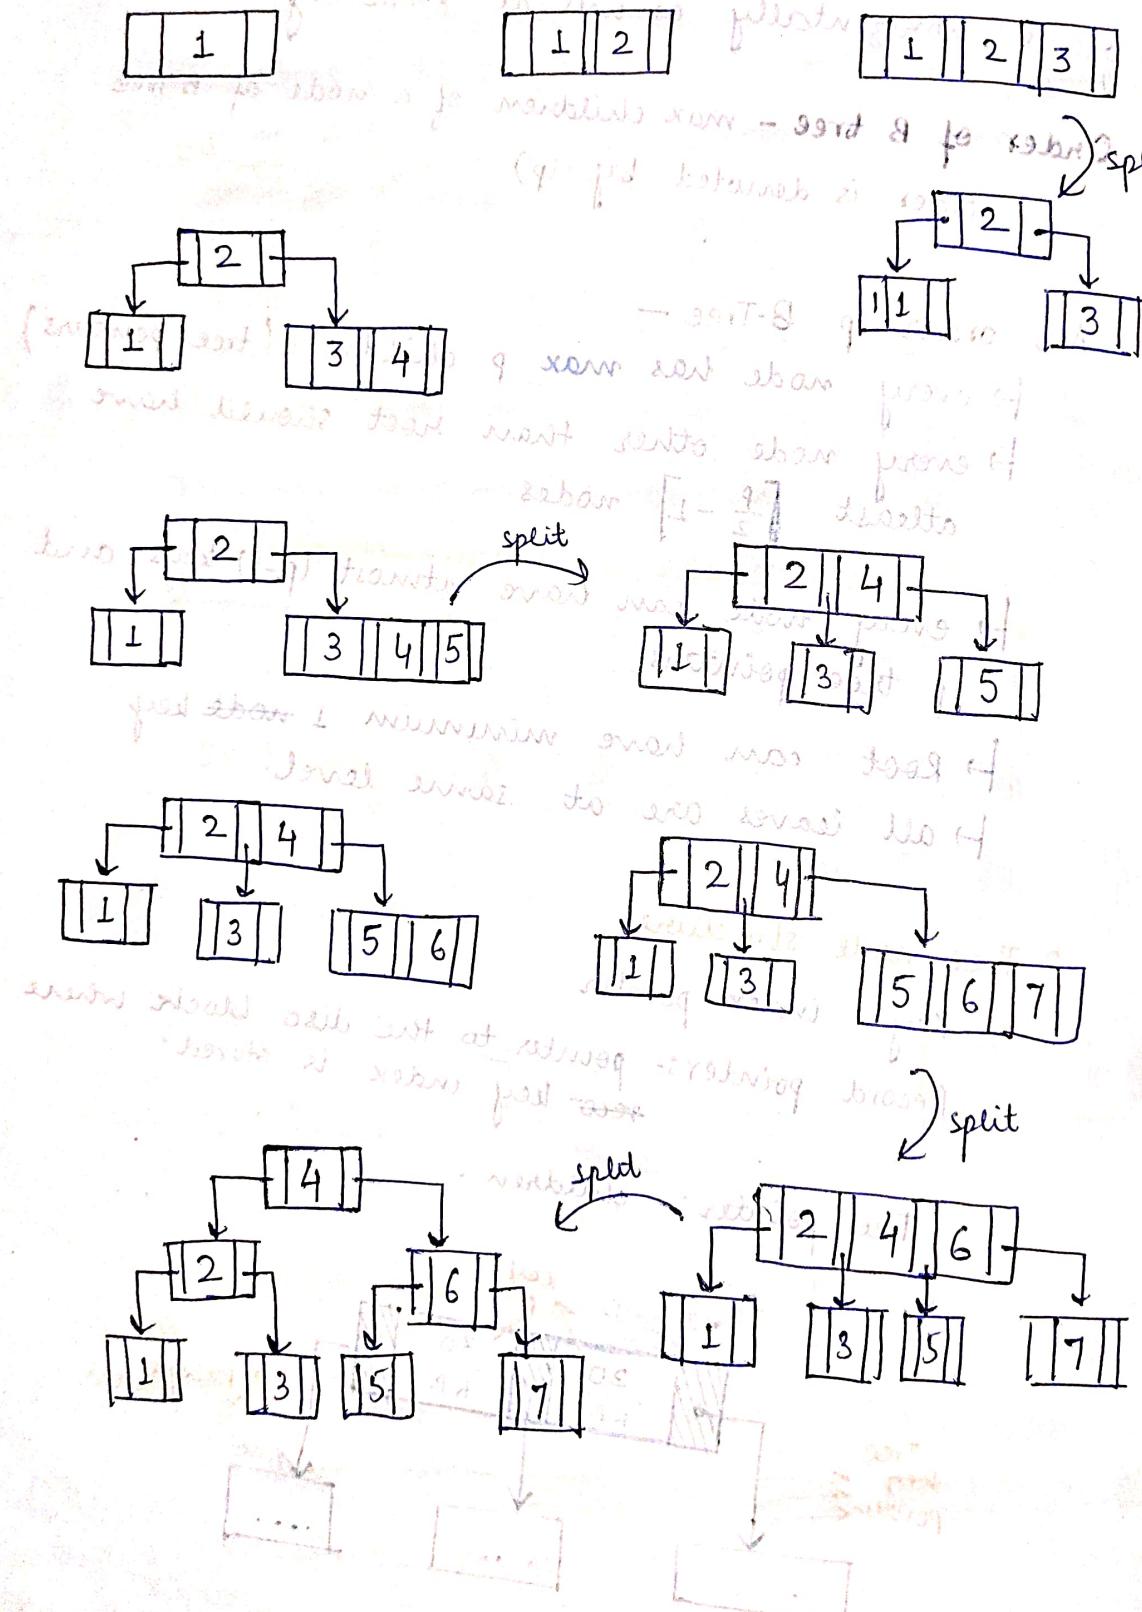
Max keys = 2

Max links (children) = 3

Min keys (except root) = $\lceil \frac{3}{2} \rceil - 1 = 1$

Max keys (except root) = $\lceil \frac{3}{2} \rceil = 2$

Insert keys 1, 2, 3, 4, 5, 6, 7



Each node of a B Tree is stored in a disk block.

If p is order of tree, it can expand on many

$$\text{max No. of keys} = p-1$$

$$\text{max. No. of record pointers} = p-1$$

$$\text{max. No. of tree pointers} = p$$

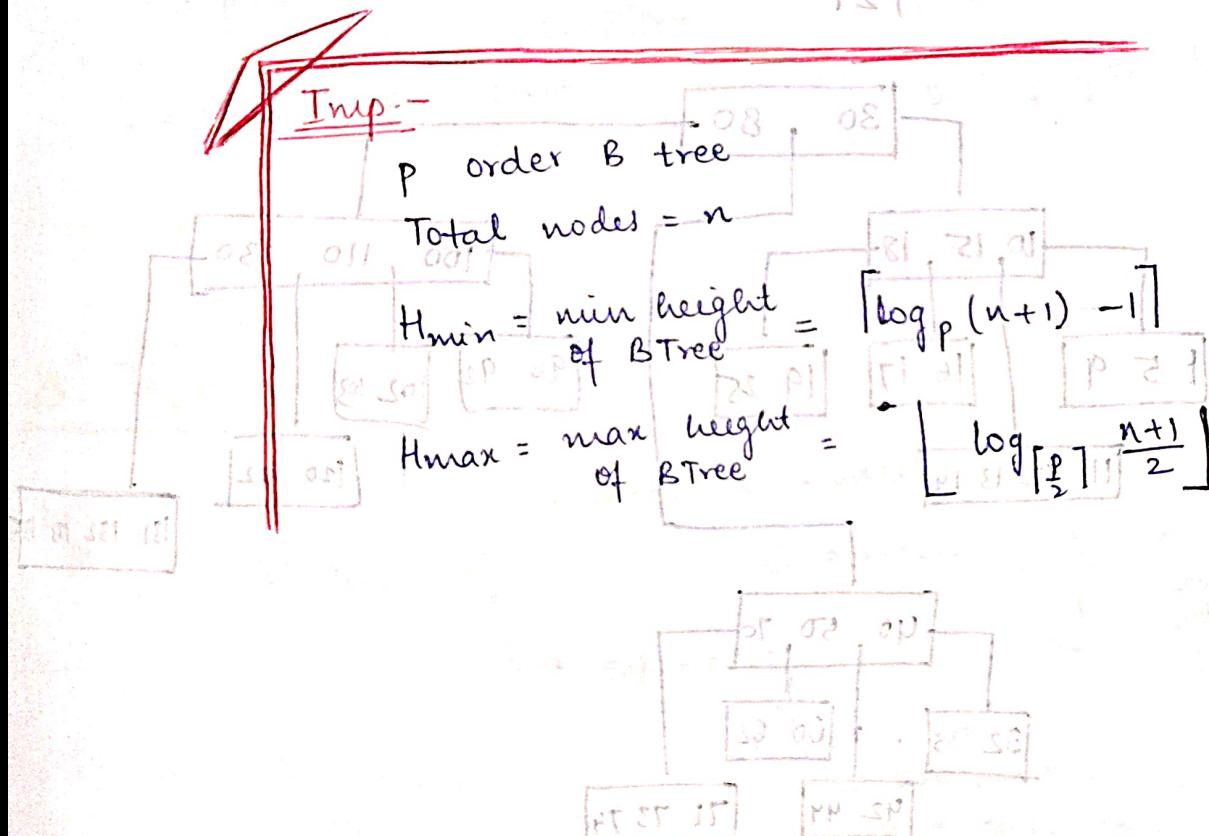
$$(p-1)(\text{key size}) + (p-1)(\text{record pointer size}) + p(\text{tree pointer size}) \leq \text{Block size}$$

formula used to find

max. value of ' p '

If block size is given

or a min block size if p is given

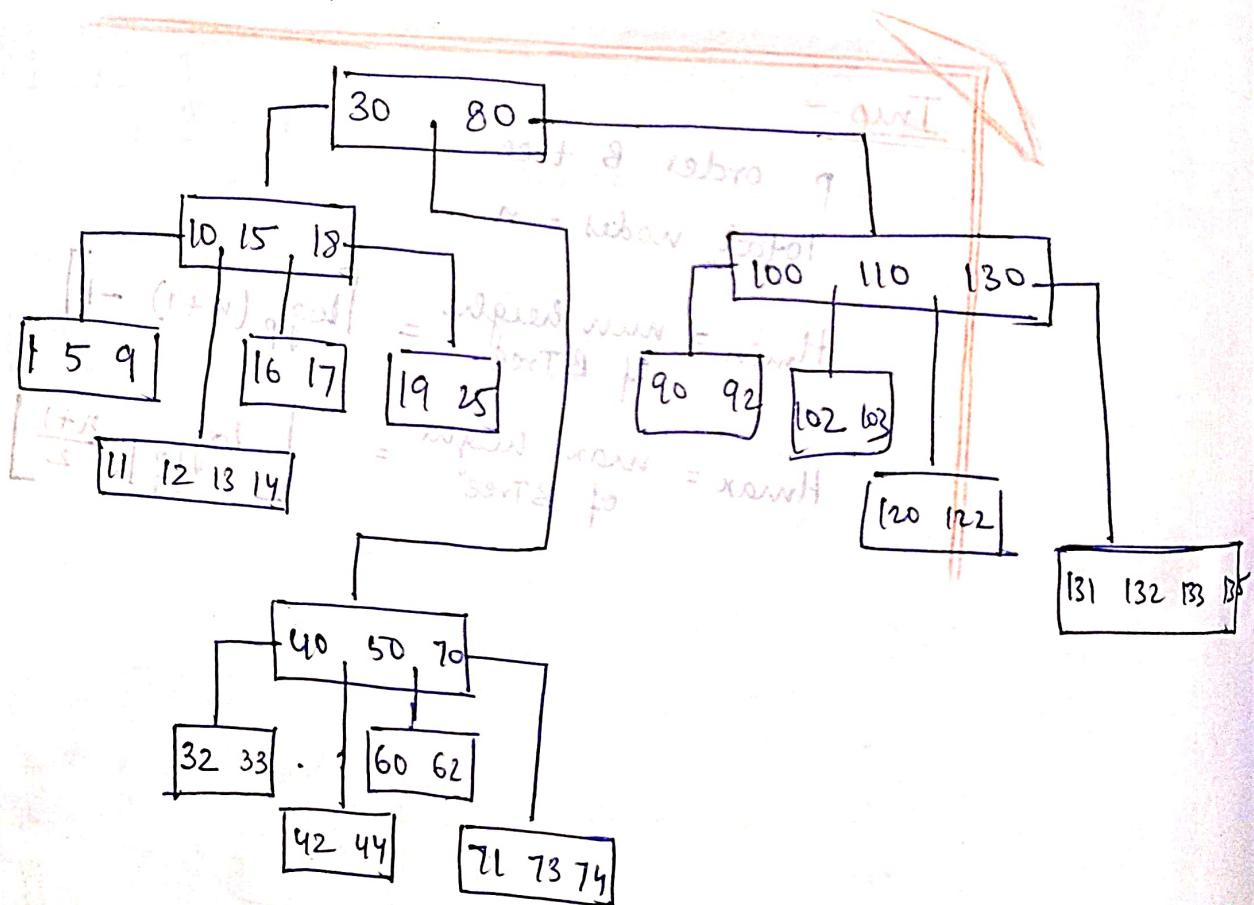


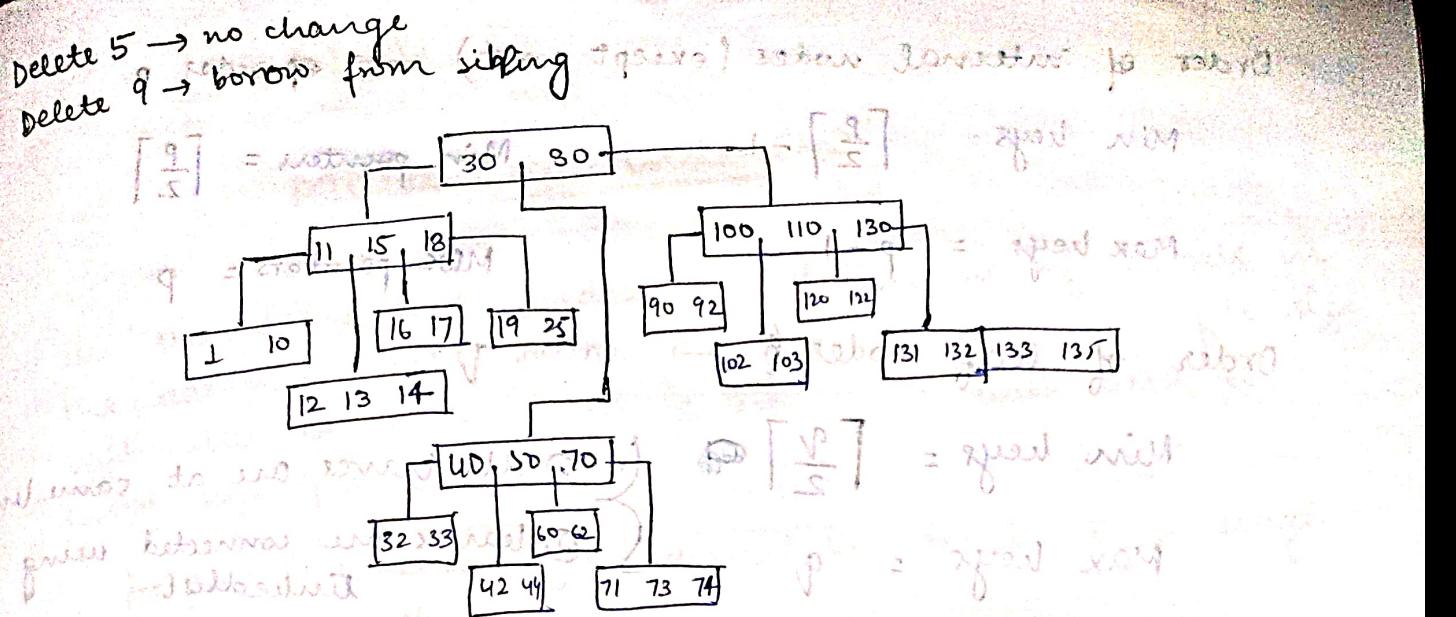
- B-Tree deletion in leaf node -
- After deletion, if no violation of new minimum keys then no changes in tree.
 - If violation of min keys, then, borrow key from sibling (rotation through parent)
 - If borrow through parent can't be possible then, merge the node with sibling & pull down the set anchor ~~node~~ key from parent.

• Btree of order 5

$$\text{max keys} = 4$$

$$\text{min keys} = \left\lceil \frac{5}{2} \right\rceil - 1 = 2$$





\Rightarrow B-tree deletion from internal node -
→ Replace the deleted value with inorder successor or inorder predecessor.
→ Now follow rules of deletion of key from leaf node.

B+ Tree

$8qas \geq (1-q) + 8(1-q)$

$8qas \geq q + 8 - 8q$

$8qas \geq 8 + 8qas \geq 8qas$

$8qas = 8 + 8qas$

$8qas = 8(1-p)$

$8qas \geq 8(1-p)$

Internal nodes
Keys
Tree pointers

leaf nodes
Keys
Record pointers
Linked list pointer

order of leaf node and internal node is different

Leaf nodes are connected to each other or (tree) or linked list

Range queries work faster in B+ trees.

Order of internal nodes (except root) \rightarrow op order P.

$$\text{Min keys} = \left\lceil \frac{P}{2} \right\rceil - 1 \quad \text{Min pointers} = \left\lceil \frac{P}{2} \right\rceil$$

$$\text{Max keys} = P - 1 \quad \text{Max pointers} = P$$

Order of leaf nodes \rightarrow order q

$$\text{Min keys} = \left\lceil \frac{q}{2} \right\rceil \quad \begin{array}{l} \text{① all leaves are at same level} \\ \text{② leaves are connected using linkedlist.} \end{array}$$

$$\text{Max keys} = q$$

- * Splitting at internal node is same as that of B Tree
- * Splitting at leaf node - middle key is duplicated & is part of right subtree.

① One node of B+ tree must be present in one disk block.

$$\text{order of B+ tree} = p$$

$$\text{Key} = 8 \text{ bytes}$$

$$\text{Block pointer} = 32 \text{ bytes}$$

$$\text{Record pointer} = 48 \text{ bytes}$$

$$\text{Disk block size} = 2048 \text{ bytes}$$

For internal node,

$$\max (p-1) \text{ keys} + p \text{ block pointers}$$

$$\therefore (p-1)8 + p(32) \leq 2048$$

$$\Rightarrow 8p - 8 + 32p \leq 2048$$

$$\Rightarrow 40p \leq 2048 + 8 \leq 2056$$

$$\Rightarrow p \leq \frac{2046}{40} = 51.125$$

For external node,

$$\max (q-1) \text{ keys}$$

$$+ (q-1) \text{ record pointers}$$

$$\therefore p = 51 \text{ (max)}$$

$$\therefore (q-1)(8) + (q-1)(48) \leq 2048$$

$$\Rightarrow (q-1)(56) \leq 2048$$

$$\Rightarrow q-1 \leq \frac{2048}{56}$$

$$\Rightarrow q-1 \leq 36.57 \Rightarrow q \leq 37.57$$

$$\therefore q = 37 \text{ (max)}$$

max order of leaf node of B+ tree = 37.

Multip

Multi-level indexing

- In secondary key indexing, if the no. of records in the database table is huge, then, the index file size will also be very huge & searching will need to access multiple disk blocks.
- entire index file cannot be stored in a single disk block.

Theory of

→ ∴ Do indexing on the index file.

< Primary indexing >

No. of blocks accessed needed to search a record is minimized.

